

# Every Face Has a Story

Automated die studies at civilisational scale:  
estimating the monetary output of the Roman Empire  
from millions of unprovenanced coin images

## 1. Abstract

## 2. Introduction

## 3. Related Work

### 3.1. Die Studies in Numismatics

### 3.2. Deep Metric Learning for Cultural Heritage

## 4. Data

### 4.1. Training Data: Schaefer / RRD

Richard Schaefer’s Roman Republican Die Project (RRDP) provides our ground-truth labels: 77587 catalogued specimen-faces, spanning 16967 hand-identified dies across 445 RRC issues (7941 obverse, 9030 reverse). We have downloaded 70371 single-face crops (median 464×455 px; 10th-percentile short side 409 px) and trained the embedding on the bulk of this corpus.

The die-size distribution is the regime that governs output estimation: 34.5% of obverse dies are seen exactly once and 50.2% at most twice, while the largest die is represented by 72 coins. For benchmarking we use the 89 obverse issues with at least 50 coins and 15 dies (34175 coins). Near-duplicate imagery is handled by the dedup pass.

### 4.2. Deployment Data: sixbid archive

The deployment corpus is the sixbid auction archive (12.4M lots). A high-recall prefilter keeps 2.58M lots with any ancient/Roman signal; of these, 601528 lots cite a catalogue and were downloaded as full-resolution two-face photographs. Because no die study can proceed without knowing which coinage a coin belongs to, every lot is first assigned a **catalogue type**: a CRRO (RRC), OCRE (RIC) or RPC identifier. Citations in the lot description resolve 459958 lots to a type; those labels train an image model (shared DINOv2 ViT-B over obverse and reverse, cosine classifier, nearest-neighbour assignment against the labelled gallery) which reaches 0.94 / 0.823 / 0.646 held-out exact-type accuracy on RRC / RIC / RPC (top-5: 0.964 / 0.915 / 0.754), and refuses non-Roman coins (Greek, Byzantine, Celtic, world) at a false-annotation rate below 1%. Applied to the uncited remainder of the candidate pool it annotates a further 355949 lots, for 930759 typed coins in total: 175303 RRC (1543 types), 621147 RIC (20542 types) and 134309 RPC (13574 types). Figure 1 summarises the corpus. The type layer supplies the population each die study is run within, the dated denominator for every issue, and the structural match between the deployment pools and the Schaefer benchmark pools.




Figure 1: The annotated corpus. (A) Dated coins by 5-year bin and metal; each coin is spread uniformly over its type's catalogue date range (interval data, no lot midpoints). RRC and RIC dates come from CRRO/OCRE; RPC Online publishes no bulk data, so its type records (province, city, reign, dating, metal) were collected page by page, with reign-level dates where a type has no explicit dating. (B) Coins per catalogue type: RRC is head-heavy (median 22 coins per type), RIC and RPC are long-tailed (medians 7 and 3). (C) Share of each catalogue's types with at least 1, 10 and 100 photographed coins. (D) From archive lots to typed coins, by annotation source. "Coins" are auction lots; resales are not yet collapsed at this level.

### 4.3. Where the Money Was Made

Every type also carries its mint (RRC and RIC as *nomisma* mint identifiers; RPC as city and province), so each annotated coin has a place as well as a date. Figure 2 follows the coin-weighted centre of gravity of state coinage in 25-year steps: fixed at Rome through the Republic, briefly pulled west by the Augustan mints of Lugdunum and Spain, back to Rome for two centuries, then moving decisively to the Danube and Balkan mints in the third-century crisis and on to Constantinople and the eastern mints after 400. The drift is the same whether periods are weighted by photographed coins or by catalogue types, so it is not an artefact of what the market sells. Mint breadth tells the same story: from one or two active state mints under the Principate to 20 in the AD 275 bin, with coins struck on average 910 km from the centroid at the peak. 6.1% of state coins and 20.7% of provincial coins belong to types with an uncertain or regional mint and are left out of the centroid.




Figure 2: The centre of gravity of Roman coinage, 225 BC – AD 500. (A) Coin-weighted centroid of state-coinage (RRC + RIC) mint locations per 25-year bin, coloured by time; grey dots are mints, sized by coins; the green path includes provincial (RPC) coinage. (B) Centroid longitude over time under three weightings: photographed coins, catalogue types, and all coinage including provincial. (C) State mints active per bin and the coin-weighted mean distance of minting from the centroid. Mint coordinates from nomisma.org; each coin is spread uniformly over its type's date range.

## 5. Method

### 5.1. Segmentation

### 5.2. Embedding: DINOv3 + LoRA

### 5.3. Nearest-Neighbour Search

### 5.4. Die Clustering

### 5.5. Output Estimation

Given a die clustering, we estimate the number of dies originally used with coverage estimators that correct for never-sampled dies. Good–Turing coverage is  $c = 1 - \frac{f_1}{n}$  ( $f_1$  = dies seen once); Chao1 gives a lower bound  $\hat{D} = d + \frac{f_1^2}{2f_2}$ . As a sanity check on Schaefer’s own labels, the obverse corpus has coverage 0.9278 over 7941 observed dies, so Chao1 implies 10961 obverse dies were used in total. Crucially there is no absolute ground-truth die count: we validate the estimator by rarefaction, not against a known  $D$ . The output figure is  $\hat{D} \times \text{coins-per-die}$ , where the multiplier — not the algorithm — is the dominant uncertainty.

### 5.6. Counting Dies Without Clustering

Cluster counting fails in one specific way: a single false link merges two dies, and enough false links percolate into giant clusters, so single-linkage is only safe at an edge precision of 0.999. The quantity we want can instead be read off pairs. Within a type whose dies carry shares  $p_i$ , two coins drawn at random share a die with probability  $q = \sum p_i^2$ , and  $D_{\text{eff}} = \frac{1}{q}$  is the effective number of dies, the same object as the effective population size in genetics. Estimating  $q$  needs only the fraction of within-type pairs the verifier accepts, corrected by its recall and false-positive rate measured on Schaefer (leave-one-type-out); a false link is then a linear bias, never a merge. Figure 3 shows the calibration and the result: on the 10 bench types with at least 50 coins, the pairwise estimate sits at 0.917 of the truth with cv 0.238, flat across thresholds, against 1.648 (cv 0.217) for single-linkage at the same threshold. The residual spread is the verifier’s per-type recall, which is the remaining lever.




Figure 3: Counting dies without clustering. (A) Within one type, a same-die pair, a different-die pair the verifier accepts at the operating point, and a typical different-die pair, with fused scores. (B) Within-type pair scores on the Schaefer bench; at threshold 120 recall is 0.586 and 0.00244 of candidate pairs are false positives. (C) Pairwise  $D_{\text{eff}}$  estimate against Schaefer truth per type (leave-one-type-out calibration), with the single-linkage cluster count at the same threshold. (D) From dies to output:  $D_{\text{eff}}$  times the conventional 10000–30000 coins per obverse die, per type; the coins-per-die assumption dominates the uncertainty.

## 6. Experiments

### 6.1. Validation Levels

## 7. Results

On the frozen RRDP test split, geometric verification fused with the DINOv3-LoRA embedding separates die-mates from hard (same-type) negatives with verification AUC 0.9591 (obverse) and 0.9476 (reverse). Pairwise quality is high; converting trustworthy links into the correct die *count* — robust to the percolation cliff — remains the binding constraint (see Discussion).

## 8. Discussion

## 9. Conclusion

## Bibliography