

Every Face Has a Story

Automated die studies at civilisational scale:
estimating the monetary output of the Roman Empire
from millions of unprovenanced coin images

1. Abstract

2. Introduction

3. Related Work

3.1. Die Studies in Numismatics

3.2. Deep Metric Learning for Cultural Heritage

4. Data

4.1. Training Data: Schaefer / RRD

Richard Schaefer’s Roman Republican Die Project (RRDP) provides our ground-truth labels: 77587 catalogued specimen-faces, spanning 16967 hand-identified dies across 445 RRC issues (7941 obverse, 9030 reverse). We have downloaded 70371 single-face crops (median 464×455 px; 10th-percentile short side 409 px) and trained the embedding on the bulk of this corpus.

The die-size distribution is the regime that governs output estimation: 34.5% of obverse dies are seen exactly once and 50.2% at most twice, while the largest die is represented by 72 coins. For benchmarking we use the 89 obverse issues with at least 50 coins and 15 dies (34175 coins). Near-duplicate imagery is handled by the dedup pass.

4.2. Deployment Data: sixbid archive

The deployment corpus is the sixbid auction archive (12.4M lots). A high-recall prefilter keeps 2.58M lots with any ancient/Roman signal; of these, 601528 lots cite a catalogue and were downloaded as full-resolution two-face photographs. Because no die study can proceed without knowing which coinage a coin belongs to, every lot is first assigned a **catalogue type**: a CRRO (RRC), OCRE (RIC) or RPC identifier. Citations in the lot description resolve 459958 lots to a type; those labels train an image model (shared DINOv2 ViT-B over obverse and reverse, cosine classifier, nearest-neighbour assignment against the labelled gallery) which reaches 0.94 / 0.823 / 0.646 held-out exact-type accuracy on RRC / RIC / RPC (top-5: 0.964 / 0.915 / 0.754), and refuses non-Roman coins (Greek, Byzantine, Celtic, world) at a false-annotation rate below 1%. Applied to the uncited remainder of the candidate pool it annotates a further 355949 lots, for 930759 typed coins in total: 175303 RRC (1543 types), 621147 RIC (20542 types) and 134309 RPC (13574 types). Figure 1 summarises the corpus. The type layer supplies the population each die study is run within, the dated denominator for every issue, and the structural match between the deployment pools and the Schaefer benchmark pools.




Figure 1: The annotated corpus. (A) Dated coins by 5-year bin and metal; each coin is spread uniformly over its type’s catalogue date range (interval data, no lot midpoints). RRC and RIC dates come from CRRO/OCRE; RPC Online publishes no bulk data, so its type records (province, city, reign, dating, metal) were collected page by page, with reign-level dates where a type has no explicit dating. (B) Coins per catalogue type: RRC is head-heavy (median 22 coins per type), RIC and RPC are long-tailed (medians 7 and 3). (C) Share of each catalogue’s types with at least 1, 10 and 100 photographed coins. (D) From archive lots to typed coins, by annotation source. “Coins” are auction lots; resales are not yet collapsed at this level.

5. Method

5.1. Segmentation

5.2. Embedding: DINOv3 + LoRA

5.3. Nearest-Neighbour Search

5.4. Die Clustering

5.5. Output Estimation

Given a die clustering, we estimate the number of dies originally used with coverage estimators that correct for never-sampled dies. Good–Turing coverage is $c = 1 - \frac{f_1}{n}$ (f_1 = dies seen once); Chao1 gives a lower bound $\hat{D} = d + \frac{f_1^2}{2f_2}$. As a sanity check on Schaefer’s own labels, the obverse corpus has coverage 0.9278 over 7941 observed dies, so Chao1 implies 10961 obverse dies were used in total. Crucially there is no absolute ground-truth die count: we validate the estimator by rarefaction, not against a known D . The output figure is $\hat{D} \times \text{coins-per-die}$, where the multiplier — not the algorithm — is the dominant uncertainty.

6. Experiments

6.1. Validation Levels

7. Results

On the frozen RRDP test split, geometric verification fused with the DINOv3-LoRA embedding separates die-mates from hard (same-type) negatives with verification AUC 0.9591 (obverse) and 0.9476 (reverse). Pairwise quality is high; converting trustworthy links into the correct die *count* — robust to the percolation cliff — remains the binding constraint (see Discussion).

8. Discussion

9. Conclusion

Bibliography